Paper Review: Reasoning and Retrieval for Complex Semi-structured Tables via Reinforced Relational Data Transformation

Setareh Babajani

1. Summary and Evaluation

1.1. Problem Definition

The paper focuses on the challenge of enabling language models to reason and retrieve information from semi-structured tables with irregular layouts such as merged cells, nested headers, and inconsistent formatting. These irregularities obscure relational structures and reduce accuracy in tasks like table question answering and fact verification. This is a practical and relevant issue since real-world tables in reports and online documents are rarely standardized. The authors propose a framework that transforms such complex tables into relational forms while preserving their semantics, improving the ability of language models to perform accurate and reliable reasoning on real-world tabular data.

1.2. Related Work

The paper builds on three lines of research, namely table transformation, step-by-step table reasoning, and reinforcement learning. Previous transformation methods such as FlashRelate or AutoTables rely on fixed rules and cannot generalize well to diverse layouts. Step-by-step reasoning methods like Chain-of-Table or ReAcTable improve interpretability but still assume clean table inputs. Reinforcement learning has been used for reasoning improvement but not for transformation. The proposed work combines these directions by using reinforcement learning to guide flexible soft operators that can adapt to complex table structures.

1.3. Approach and Methods

The proposed model, called TabFormer, introduces a chain of semantic transformations executed through soft operators. These operators perform contextual modifications such as pivoting, unmerging, and restructuring while preserving semantic consistency. The output is expressed in SQL format, which allows reasoning across tables using queries. The model is trained in two phases: (1) Supervised warm-up training which uses synthetic data generated by AutoTables to provide the model with initial table transformation capabilities, (2) Reinforcement learning fine-tuning which applies the Proximal Policy Optimization algorithm with two main rewards, one measuring relational normalization and another measuring reasoning correctness. A reranking and fallback mechanism is also applied to maintain stability when transformations fail.

1.4. Experiments and Evaluation

The authors evaluate their framework on four public benchmarks including WikiTQ, HiTab, MultiHiertt, and TabFact. TabFormer achieves consistent improvements ranging from 4% to

17% compared to recent baselines such as TableLlama, ReAcTable, and EEDP. The ablation results confirm that both the reward structure and the transformation chain contribute meaningfully to performance. The experiments are comprehensive, covering multiple reasoning types and evaluation metrics. However, the paper could benefit from more discussion on computational efficiency and scalability when dealing with larger or more complex tables.

1.5. Contributions and Impact

The paper makes three notable contributions. First, it presents a reinforcement learning framework for converting semi-structured tables into relational form using semantic soft operators. Second, it introduces a dual reward mechanism that improves both transformation quality and reasoning accuracy without relying on labeled data. Third, it provides experimental evidence showing that reasoning over normalized tables improves accuracy significantly. These contributions are valuable for future developments in structured data understanding and retrieval-augmented reasoning systems.

2. Critical Analysis

2.1. Weaknesses and Limitations

The framework is effective but has some limitations. It mainly handles reasoning within a single document and does not support multi-document retrieval, which limits its use in larger systems. Experiments focus on small to medium tables, leaving performance on very large tables uncertain. The relational normalization reward depends on GPT-4o, introducing reliance on a proprietary model, and the stepwise reasoning reward requires multiple language models, increasing computational cost. The synthetic warm-up data lacks realistic artifacts such as merged or irregular cells, making the model heavily dependent on reinforcement learning for generalization. These issues highlight opportunities to improve scalability, independence from external models, and efficiency.

2.2. Clarity and Readability

The paper is well organized and mostly clear. The structure follows a logical flow from problem definition to proposed method and experimental validation. Figures and tables effectively summarize the key results and workflow. The examples of transformed tables are particularly helpful for understanding the soft operator mechanism. However, the section describing reinforcement learning introduces several new notations that could be explained more clearly. A small illustrative example showing how rewards are computed during training would make the learning process easier to follow.

2.3. Possible Improvements

Future work could expand the system to handle large-scale and multi-document reasoning tasks. Introducing compression or chunking techniques would allow the model to process larger tables efficiently. Developing a normalization metric that does not depend on external models such as GPT-4o could improve independence and reproducibility. Training a smaller learned reward model might reduce computational cost while keeping performance stable. A deeper qualitative study of soft operators on challenging table types would also strengthen the argument for their generalization capability. Finally, integrating TabFormer into retrieval-augmented systems could extend its practical usefulness in broader reasoning pipelines.